

Software

Introduction to computer programming

Management and Artificial Intelligence



How can a computer execute a program?



The **CPU** is the responsible for executing the program code:

- It interprets the instructions one by one from the **RAM**
- The instructions must be written in **machine code**
- The machine code is:
 - **Low level**: relatively simple operations
 - **CPU specific**: different CPUs have different machine codes
 - **Binary format**: humans struggle to read it

Machine code: Intel/AMD vs ARM

A simple sequence of instructions to compute the sum of two integer numbers would look like this in machine code:

Intel/AMD 64 bit machine code

```
f3 0f 1e fa
55
48 89 e5
89 7d fc
89 75 f8
8b 55 fc
8b 45 f8
01 d0
5d
c3
```

ARM 64-bit machine code

```
ff 43 00 d1
e0 0f 00 b9
e8 0f 40 b9
e9 0f 40 b9
00 7d 09 1b
ff 43 00 91
c0 03 5f d6
```

- Each instruction is stored into its own line for clarity, but in practice, they would be stored contiguously in memory.
- The machine code is stored in a binary format, while in this example we have used hexadecimal notation for clarity.
- Intel and AMD processors use the same machine code format (with only minor differences).
- Other processors (e.g., ARM used by Apple, Raspberry Pi, Huawei, etc.) use a completely different machine code format.

From machine code to assembly code

To make machine code understandable for human, we can translate it to assembly code:

- **Low level** [same as machine code]: relatively simple operations
- **CPU specific** [same as machine code]: different CPUs have different machine codes
- **Human readable** [differently from machine code]: CPU struggles to read it!

Intel/AMD 64 bit assembly

```
endbr64
push    %rbp
mov     %rsp,%rbp
mov     %edi,-0x4(%rbp)
mov     %esi,-0x8(%rbp)
mov     -0x4(%rbp),%edx
mov     -0x8(%rbp),%eax
add     %edx,%eax
pop     %rbp
ret
```

ARM 64-bit assembly

```
sub     sp, sp, #16
str     w0, [sp, #12]
ldr     w8, [sp, #12]
ldr     w9, [sp, #12]
mul     w0, w8, w9
add     sp, sp, #16
ret
```

Without getting into details, we can see that this code is doing some basic operations, such as addition, multiplication, copy/move of values from/to registers (which are different from CPU to CPU), etc.

NOTE: Translation from/to machine code to/from assembly code is a one-to-one non-ambiguous mapping. Translating from machine code to assembly code is called **disassembly**, while the opposite is called **assembly**.

Writing low level code is impractical

Major problems when working with machine/assembly code are:

1. Programmers have to write a lot of code to perform simple operations
2. Code is not portable across different hardware platforms
3. Programmers have to deal with hardware details
4. Programming in machine/assembly code is hard and error-prone

Yet, CPUs only understand and run machine code.

How can programmers survive?

From low-level to high-level programming

Most programmers use high-level languages, such as Python, Javascript, or PHP, to write software.

However, the CPU does not understand high-level languages. Never!

Code written in high-level languages must be translated into machine code. Two main strategies:

1. **Compilation:**


- Compilers translate high-level code into machine code at *compile time*
- Compile time means when the developer writes the code.
- Machine code is packaged into an executable.
- One executable for each platform (Operating Systems and CPUs).
- Developers keeps the high-level code, while users get the low-level code.

2. **Interpretation:**


- Interpreters translate high-level code into machine code at *running time*.
- Running time means when the user runs the code.
- The code is portable and platform independent.
- Developers distribute the high-level code, users exploits the interpreter to run the code.

Compiled code is typically more efficient than interpreted code.

Compilation



Interpretation



Why do we need to learn a programming language?

Most everyday tasks require to perform:

- repetitive...
- long sequences of...
- simple steps

This scenario is *exactly* where a computer excels!

Why do we need to learn a programming language?

Most everyday tasks require to perform:

- repetitive...
- long sequences of...
- simple steps

This scenario is *exactly* where a computer excels!

Computers are incredibly fast, accurate, and stupid. Human beings are incredibly slow, inaccurate, and brilliant. Together they are powerful beyond imagination.

-- Albert Einstein

Language: syntax and semantics

Programming languages have:

- **Syntax:** it defines which strings of characters and symbols are well formed.
- **Semantics:** it defines the meaning of the well formed strings.

The same is true even for natural languages (English, Italian, etc.):

- **Syntax:**
 - *"I are big"* is syntactically incorrect (*"I am big"* is correct).
 - *"The square circle is big"* is a well formed sentence.
- **Semantics:**
 - *"The square circle is big"* is semantically incorrect (nonsensical because "square" and "circle" contradict).

An important difference between programming languages and natural languages is that programming languages are designed to be **unambiguous**, which means that there is no room for interpretation.

Programming language paradigms

Two major computing paradigms:

- **Imperative programming:**

- we describe how the program must accomplish the task
- language statements describe the steps (how to do), just like a **recipe**
- close to the paradigm adopted by machine code

- **Declarative programming:**

- we describe what to do (define the end goal), without saying how (no step-by-step)
- language statements describe facts (what to do)
- paradigm adopted, e.g., by Database programming languages (SQL)
- distant from the paradigm adopted by machine code

Imperative programming is what we will see in this course.

NOTE: In many cases, code will be a mixture of both designs too, so it's not always black-and-white.

Solving a task → Algorithm → Recipe

An algorithm features:

- a sequence of **simple steps**
- **flow of control** to specify when each step must be executed
- a way of determining when to **stop**.

...just like a recipe!

Even a plain $1 + 2 + 3 + 4 + 5 = ?$ is an algorithm!

A slightly more complex example: Newton's method

The square root of x is a number y such that $y \cdot y = x \leftarrow$ **declarative**!

Let's see an **imperative** version:

1. Start with a guess g
2. If $g \cdot g$ is "close enough" to x :
 - A. Let $y = g$
 - B. Stop
3. Otherwise: a. Let $g \leftarrow \frac{1}{2} \left(g + \frac{x}{g} \right)$ b. Repeat from step #2

Square root of 16 via Newton's method:

g	$g \cdot g$	$\frac{1}{2} \left(g + \frac{x}{g} \right)$
3	$3 \cdot 3 = 9$	$\frac{1}{2} \left(3 + \frac{16}{3} \right) \simeq 4.17$
4.17	$4.17 \cdot 4.17 \simeq 17.39$	$\frac{1}{2} \left(4.17 + \frac{16}{4.17} \right) \simeq 4.0035$
4.0035	$4.0035 \cdot 4.0035 \simeq 16.028$	$\frac{1}{2} \left(4.0035 + \frac{16}{4.0035} \right) \simeq 4$

As already pointed out, we cannot write an algorithm in our own natural language (e.g., English). To make an algorithm interpretable by a computer, we need to learn a **programming language**.

MANY different high-level languages




There exists thousands of high-level languages:

- They have been designed with different goals in mind
- Some are general purpose, some are domain specific (e.g., web)
- Some gives more control to the user, some hide more details
- Some are compiled, some are interpreted

In this course, we will focus on a single high-level language: ***but which we should choose?***

Popular programming languages



See the full [video](#)

Python



Benefits:

- **Easy to learn:** designed to be readable, quick to learn
- **High-level language:** abstracting away most low-level aspects of a computer
- **Large community and great support:** several libraries allow to build complex projects with just a few lines of code
- **Adequate for most tasks:** it shows its limits when performance is essential and low-level details play a role. You will hardly work on anything that cannot be implemented in Python

Python:

- is an **imperative programming language**
- with minor flavors of declarative (functional) programming traits.

Computers run complex software stacks

Python is nice and powerful, but computers are expected to run **MANY** programs at the same time:

- Browser to visit web pages (e.g., Chrome)
- Office editor to write documents
- File manager to organize files
- Messaging applications (e.g., Telegram)
- Video streaming client (e.g., Netflix)
- ...and way more!

Some of these may be written in Python, but others are written in other languages.

Computers run complex software stacks

Python is nice and powerful, but computers are expected to run **MANY** programs at the same time:

- Browser to visit web pages (e.g., Chrome)
- Office editor to write documents
- File manager to organize files
- Messaging applications (e.g., Telegram)
- Video streaming client (e.g., Netflix)
- ...and way more!

Some of these may be written in Python, but others are written in other languages.

Hence, **a computer cannot just run a single Python program**. Moreover, computers are expected to:

- embrace the concept of user accounts
- provide access only after user authentication
- protect files across users and prevent hardware damage
- allow only admin users to install other programs
- handle network connections, power management

This is a **LOT** of **HARD** work: we need an **Operating System (OS)**.

Operating Systems

Modern computing systems are based on **operating systems**, which are software designed to allow the execution of applications (programs) on the available hardware.

The main tasks of an operating system are:


- **Process management** (creation, termination, CPU scheduling, etc.): any application (program) running on your OS is associated with a **process**.
- **Resource allocation** to processes (RAM, storage, network connections)
- **Time management** (clock, timers)
- **User management** for multi-user systems (access privileges to files, programs)
- **Inter-process communication** (locally or across networked systems)
- **Error handling** during process execution or system operations
- **Protection and security**, isolating processes to prevent interference with other processes, the OS, or hardware (preventing data theft or malicious/buggy operations)

Design goals of modern Operating Systems

Modern operating systems are designed with several key objectives:

- **Multitasking:** Ability to run multiple applications simultaneously
- **Preemption:** Capability to take CPU control from one process and give it to another
- **Progress & fairness:** Regular CPU allocation to all processes to ensure computational progress
- **High throughput:** Maximizing the number of instructions executed per time unit
- **Low latency:** Minimizing response time for system operations (critical in real-time environments like automotive safety systems)


Overview Operating System



From the bottom:

- The **hardware** is the physical part of the computer, containing CPU, RAM, disks and other devices.
- To program the hardware, each architecture defines the **Instruction Set Architecture (ISA)**, which defines machine code instruction syntax and semantics.
- The platform (operating system on a specific architecture, e.g., Windows x86) defines the **Application Binary Interface (ABI)**, which defines conventions (rules) on how different programs can interact, reuse code, and interface with the operating system.


Overview Operating System (cont'd)



The **operating system**:

- has full control over the hardware:
 - it runs in **supervisor (kernel) execution mode**
 - it can execute **ALL** instructions from the ISA.
- controls which applications should run on the CPU cores at any time
- plays the role of intermediary between applications and the hardware
- manages four key aspects: **processes, memory management, file and I/O, time**

Overview Operating System (cont'd)



In **user execution mode**:

- Code can run only a limited (safe) subset of ISA
- Privileged instructions must be delegated to the OS following ABI
- Processes can be related to:
 - built-in applications provided by the OS: e.g., file manager
 - applications installed by the users: e.g., Chrome (web browser by Google)
- Applications use system libraries (subprograms for standard tasks, e.g., reading a file, make a network communication) to avoid reinventing the wheel when not needed.

CPU Cores and Operating System Multitasking

- CPUs traditionally had only a single core (and many still have few cores)
- Yet we can run many applications simultaneously

How is this possible?

CPU Cores and Operating System Multitasking

- CPUs traditionally had only a single core (and many still have few cores)
- Yet we can run many applications simultaneously

How is this possible?

The Operating System creates the **illusion of parallel execution**:

- CPU time is divided into small time slots (milliseconds)
- Each application gets a turn to use the CPU
- When a time slot expires, the OS performs *context switching*
- Applications take turns using the CPU so quickly that users perceive them running simultaneously

This technique is called **time-sharing** or **preemptive multitasking** and it is the essential trait of all modern operating systems.

CPU Cores and Operating System Multitasking

- CPUs traditionally had only a single core (and many still have few cores)
- Yet we can run many applications simultaneously

How is this possible?

The Operating System creates the **illusion of parallel execution**:

- CPU time is divided into small time slots (milliseconds)
- Each application gets a turn to use the CPU
- When a time slot expires, the OS performs *context switching*
- Applications take turns using the CPU so quickly that users perceive them running simultaneously

This technique is called **time-sharing** or **preemptive multitasking** and it is the essential trait of all modern operating systems.

If you have:

- 1 core with preemptive multitasking: illusion of parallel execution of an infinite number of apps
- k cores without preemptive multitasking: up to k applications executing in parallel
- k cores with preemptive multitasking: illusion of parallel execution of an infinite number of applications, with up to k actually executing in parallel at any given time

Popular Operating Systems

Windows

- Proprietary, closed-source OS developed by Microsoft for:
 - Intel/AMD CPUs
 - ARM (still early stage)
- Most widely used desktop OS globally
- Known for user-friendly interface and software compatibility



macOS

- Proprietary, closed-source OS developed by Apple for:
 - PowerPC CPUs (before 2006)
 - Intel CPUs (before 2021)
 - Apple Silicon CPUs (now)
- Known for sleek design and stability
- Popular among creative professionals



GNU/Linux

- Combination of the Linux kernel and GNU userspace libraries
- Packaged by different *distributions* (Ubuntu, Fedora, Debian, Gentoo, Arch)
- **Open Source**: you can see and change its code!
- Powers most web servers and supercomputers (even Microsoft uses it in their cloud!)



Android

- Linux-based OS developed by Google
- Most popular mobile OS worldwide
- Open ecosystem with extensive app marketplace



iOS

- Proprietary, closed-source OS for Apple mobile devices
- Second most popular mobile OS worldwide
- Closed ecosystem with extensive app marketplace



Learn the OS Essentials

In the next lecture, we will learn about and work with essential operating system concepts.

